

Artificial intelligence document assistants help users query PDF content through retrieval augmented generation pipelines. Local models via Ollama enable offline embeddings and chat without cloud APIs. ChromaDB stores vector indexes per document. FastAPI exposes upload and streaming chat endpoints for developers. PyMuPDF extracts text from each page before chunking with overlap. Recursive character splitters preserve paragraph boundaries when possible. Semantic search retrieves the most relevant passages for each user question. Conversation memory keeps recent turns for contextual follow up questions. Graceful degradation returns clear errors when Ollama is not running locally.

Artific
gener
Chrom
devel
splitte
passa
quest

Artificial intelligence document assistants help users query PDF content through retrieval augmented generation pipelines. Local models via Ollama enable offline embeddings and chat without cloud APIs. ChromaDB stores vector indexes per document. FastAPI exposes upload and streaming chat endpoints for developers. PyMuPDF extracts text from each page before chunking with overlap. Recursive character splitters preserve paragraph boundaries when possible. Semantic search retrieves the most relevant passages for each user question. Conversation memory keeps recent turns for contextual follow up questions. Graceful degradation returns clear errors when Ollama is not running locally.

Artific
gener
Chrom
devel
splitte
passa
quest